

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Theory:

Container-based microservices architectures have revolutionized how development and operations teams test and deploy modern software. Containers allow companies to scale and deploy applications more efficiently, but they also introduce new challenges, adding complexity by creating a whole new infrastructure ecosystem.

Today, both large and small software companies are deploying thousands of container instances daily. Managing this level of complexity at scale requires advanced tools. Enter Kubernetes.

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications. Kubernetes has quickly become the de facto standard for container orchestration and is the flagship project of the Cloud Native Computing Foundation (CNCF), supported by major players like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat. Kubernetes simplifies the deployment and operation of applications in a microservice architecture by providing an abstraction layer over a group of hosts. This allows development teams to deploy their applications while Kubernetes takes care of key tasks, including:

- Managing resource consumption by applications or teams
- Distributing application load evenly across the infrastructure
- Automatically load balancing requests across multiple instances of an application
- Monitoring resource usage to prevent applications from exceeding resource limits and automatically restarting them if needed
- Moving application instances between hosts when resources are low or if a host fails
- Automatically utilizing additional resources when new hosts are added to the cluster
- Facilitating canary deployments and rollbacks with ease

Necessary Requirements:

- EC2 Instance: The experiment required launching a t2.medium EC2 instance with 2 CPUs, as Kubernetes demands sufficient resources for effective functioning.

● **Minimum Requirements:**

- Instance Type: t2.medium
- CPUs: 2
- Memory: Adequate for container orchestration.

This ensured that the Kubernetes cluster had the necessary resources to function smoothly

Sign in to AWS Management Console:

Go to AWS Management Console.

Log in with your account credentials.

Navigate to EC2 Service:

In the AWS Console, search for EC2 and select it to open the EC2 dashboard.


Go to "Security Groups":

On the left-hand navigation pane, under Network & Security, click on Security Groups.

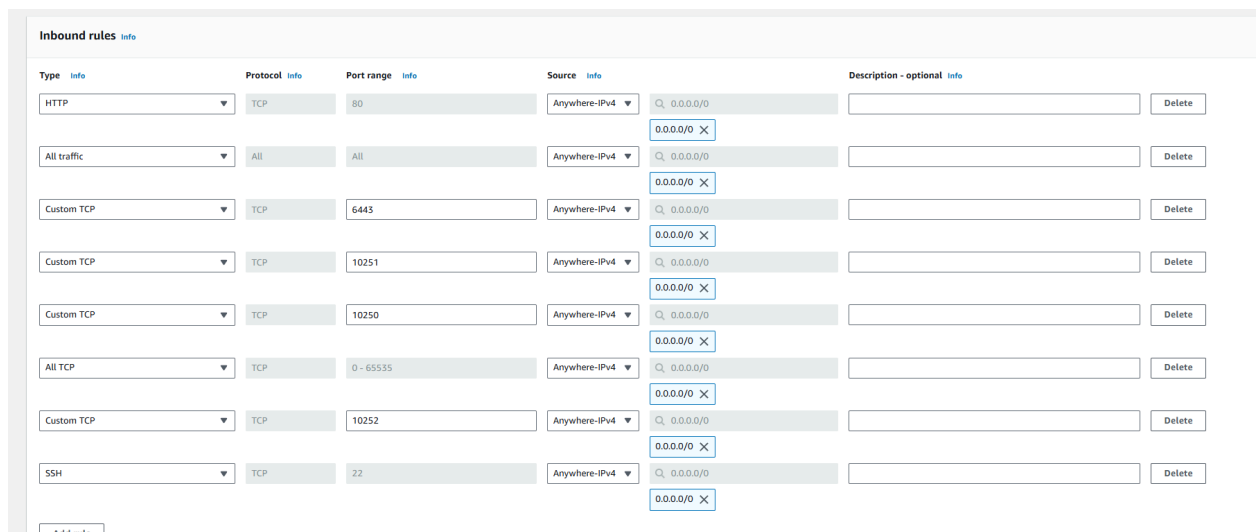
Create Security Group:

Click the Create Security Group button at the top.

Create 2 Security Groups for Master and Nodes and add the following rules inbound rules in those Groups.



MASTER:



NODE:

Inbound rules

Type

Protocol

Port range

Source

Description - optional

All traffic	All	All	Anywhere-IPv4	0.0.0.0/0		Delete
SSH	TCP	22	Anywhere-IPv4	0.0.0.0/0		Delete
Custom TCP	TCP	10250	Anywhere-IPv4	0.0.0.0/0		Delete
All TCP	TCP	0 - 65535	Anywhere-IPv4	0.0.0.0/0		Delete
Custom TCP	TCP	30000 - 32767	Anywhere-IPv4	0.0.0.0/0		Delete
HTTP	TCP	80	Anywhere-IPv4	0.0.0.0/0		Delete

Add rule

Step 1: Log in to your AWS Academy/personal account and launch 3 new Ec2 Instances. Select Ubuntu as AMI and t2.micro (because in academic account only t2.micro is present) as Instance Type and create a key of type RSA with .pem extension and move the downloaded key to the new folder. We can use 3 Different keys or 1 common key also.

Launch an instance

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags

NameMaster

Add additional tags

Application and OS Images (Amazon Machine Image)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search your full catalog including 1000s of application and OS images

RecentsQuick Start

Amazon Linux

macOS

Ubuntu

Windows

Red Hat

SUSE L

Browse more AMIs

Amazon Machine Image (AMI)

Ubuntu Server 24.04 LTS (HVM), 55D Volume Type

Free tier eligible

Number of instances

1

Software Image (AMI)

Canonical, Ubuntu, 24.04, amd64...read more

Virtual server type (instance type)

t2.micro

Firewall (security group)

-

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

CancelLaunch instanceReview commands

Instance type

Instance type

t2.micro

Family: t2 1 vCPU 1 GiB Memory Current generation: true

On-Demand Windows base pricing: 0.0162 USD per Hour

On-Demand SUSE base pricing: 0.0116 USD per Hour

On-Demand RHEL base pricing: 0.026 USD per Hour

On-Demand Linux base pricing: 0.0116 USD per Hour

Additional costs apply for AMIs with pre-installed software

All generationsCompare instance types

Key pair (login)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

exp3

Create new key pair

▼ Network settings

Info

Edit

Network

Info

vpc-0d4c0d8f48c2e4508

Subnet

Info

No preference (Default subnet in any availability zone)

Auto-assign public IP

Info

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

Common security groups

Info

Select security groups

MasterGroup sg-097fc30a345c1a537

VPC: vpc-0d4c0d8f48c2e4508

Compare security group rules

Security groups that you add or remove here will be added to or removed from all your network interfaces.

Number of instances

Info

1

Software Image (AMI)

Canonical, Ubuntu, 24.04, amd64...read more

ami-0e86e20dae9224db8

Virtual server type (instance type)

t2.micro

Firewall (security group)

MasterGroup

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per

Do Same for 2 Nodes and use security groups of Node for that.

Firewall (security groups)

Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

Common security groups

Info

Select security groups

NodeGroup sg-030c0a1b62a1e9894

VPC: vpc-0d4c0d8f48c2e4508

Compare security group rules

Security groups that you add or remove here will be added to or removed from all your network interfaces.

Firewall (security group)

NodeGroup

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable)

Step 2: After creating the instances click on Connect & connect all 3 instances and navigate to SSH Client.

aws

Services

Search

[Alt+S]

EC2 Dashboard

EC2 Global View

Events

Console-to-Code

Preview

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity

Instances (7)

Info

Find instance by attribute or tag (case-sensitive)

All states

Last updated less than a minute ago

Connect

Instance state

Actions

Launch instance

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4
<input type="checkbox"/>	Master	i-0c6758f4d6eea8fc	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d	ec2-54-208-20-141.compute-1.amazonaws.com	54.208.20.141
<input type="checkbox"/>	MASTER	i-0044a08ca3541e460	Terminated	t2.micro	-	View alarms +	us-east-1a	-	-
<input type="checkbox"/>	nagios-host	i-04709b3512d9f750f	Terminated	t2.micro	-	View alarms +	us-east-1d	-	-
<input type="checkbox"/>	node1	i-041464f92af63c03e	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d	ec2-35-170-201-119.compute-1.amazonaws.com	35.170.201.119
<input type="checkbox"/>	node2	i-0d57570c061c25ae1	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d	ec2-184-73-147-235.compute-1.amazonaws.com	184.73.147.235
<input type="checkbox"/>	SLAVE	i-00d0a5a729eb2e648	Terminated	t2.micro	-	View alarms +	us-east-1a	-	-
<input type="checkbox"/>	SLAVE2	i-0426944f0d33bd940	Terminated	t2.micro	-	View alarms +	us-east-1a	-	-

Step 3: Now open the folder in the terminal 3 times for Master, Node1 & Node 2 where our .pem key is stored and paste the Example command (starting with ssh -i) in the terminal.
ssh -i "<PATH TO FILE>exp3.pem" ubuntu@ec2-54-208-20-141.compute-1.amazonaws.com

MASTER:

CONNECT TO INSTANCE [info](#)

Connect to your instance i-Oc67658f4d6eea8fc (Master) using any of these options

EC2 Instance Connect	Session Manager	SSH client	EC2 serial console
----------------------	-----------------	-------------------	--------------------

Instance ID

i-Oc67658f4d6eea8fc (Master)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is exp3.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "exp3.pem"`
4. Connect to your instance using its Public DNS:
 `ec2-54-208-20-141.compute-1.amazonaws.com`

Example:

`ssh -i "exp3.pem" ubuntu@ec2-54-208-20-141.compute-1.amazonaws.com`

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

NODE 1&2:

EC2 Instance Connect	Session Manager	SSH client	EC2 serial console
----------------------	-----------------	-------------------	--------------------

Instance ID

i-0414d4f92af63c03e (node1)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is exp3.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "exp3.pem"`
4. Connect to your instance using its Public DNS:
 `ec2-35-170-201-119.compute-1.amazonaws.com`

Example:

`ssh -i "exp3.pem" ubuntu@ec2-35-170-201-119.compute-1.amazonaws.com`

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

EC2 Instance Connect	Session Manager	SSH client	EC2 serial console
----------------------	-----------------	-------------------	--------------------

Instance ID

i-0d57570c061c25ae1 (node2)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is exp3.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "exp3.pem"`
4. Connect to your instance using its Public DNS:
 `ec2-184-73-147-235.compute-1.amazonaws.com`

Example:

`ssh -i "exp3.pem" ubuntu@ec2-184-73-147-235.compute-1.amazonaws.com`

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

```
ubuntu@ip-172-31-82-192: ~  
s.  
This key is not known by any other names.  
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes  
Warning: Permanently added 'ec2-54-208-20-141.compute-1.amazonaws.com' (ED25  
519) to the list of known hosts.  
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)  
  
* Documentation:  https://help.ubuntu.com  
* Management:    https://landscape.canonical.com  
* Support:       https://ubuntu.com/pro  
  
System information as of Thu Sep 26 15:21:26 UTC 2024  
  
System load:  0.0      Processes:            104  
Usage of /:   22.7% of 6.71GB   Users logged in:     0  
Memory usage: 20%      IPv4 address for enX0: 172.31.82.192  
Swap usage:   0%  
  
Expanded Security Maintenance for Applications is not enabled.  
  
0 updates can be applied immediately.  
  
Enable ESM Apps to receive additional future security updates.  
See https://ubuntu.com/esm or run: sudo pro status  
  
The list of available updates is more than a week old.  
To check for new updates run: sudo apt update  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
ubuntu@ip-172-31-82-192:~$
```

```
ubuntu@ip-172-31-86-84: ~  
Enable ESM Apps to receive additional future security updates.  
See https://ubuntu.com/esm or run: sudo pro status  
  
The list of available updates is more than a week old.  
To check for new updates run: sudo apt update  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
ubuntu@ip-172-31-86-84:~$
```

```
ubuntu@ip-172-31-86-8: ~  
Enable ESM Apps to receive additional future security updates.  
See https://ubuntu.com/esm or run: sudo pro status  
  
The list of available updates is more than a week old.  
To check for new updates run: sudo apt update  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
ubuntu@ip-172-31-86-8:~$
```

Step 4: Run on Master, Node 1, and Node 2 the below commands to install and setup Docker in Master, Node1, and Node2.

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
```

```
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
```

```
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
```

```
ubuntu@ip-172-31-82-192:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
```

```
ubuntu@ip-172-31-82-192:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/null
ubuntu@ip-172-31-82-192:~$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
Adding repository.
Press [ENTER] to continue or Ctrl-c to cancel.
Adding deb entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Adding disabled deb-src entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:4 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:5 https://download.docker.com/linux/ubuntu noble InRelease [48.8 kB]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Packages [15.3 kB]
Get:7 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Components [212 B]
Get:8 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 c-n-f Metadata [116 B]
Get:9 https://download.docker.com/linux/ubuntu noble/stable amd64 Packages [15.3 kB]
Fetched 29.1 MB in 6s (4873 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
ubuntu@ip-172-31-82-192:~$
```

sudo apt-get update

sudo apt-get install -y docker-ce

```
ubuntu@ip-172-31-82-192:~$ sudo apt-get update
sudo apt-get install -y docker-ce
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  containerd.io docker-buildx-plugin docker-ce-cli
  docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
Suggested packages:
  aufs-tools cgroupfs-mount | cgroup-lite
The following NEW packages will be installed:
  containerd.io docker-buildx-plugin docker-ce docker-ce-cli
  docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
Processing triggers for man-db (2.12.0-4build2) ...
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-82-192:~$ |
```

sudo mkdir -p /etc/docker

cat <<EOF | sudo tee /etc/docker/daemon.json

```
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
```

```
No VM guests are running outdated hypervisor (qemu) binaries on this host
ubuntu@ip-172-31-82-192:~$ sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-82-192:~$
```

sudo systemctl enable docker

sudo systemctl daemon-reload

sudo systemctl restart docker

```
ubuntu@ip-172-31-82-192:~$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with /usr/lib
/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-82-192:~$
```

Step 5: Run the below command to install Kubernetes.

curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o

/etc/apt/keyrings/kubernetes-apt-keyring.gpg

echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]

https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee

/etc/apt/sources.list.d/kubernetes.list

```
ubuntu@ip-172-31-82-192:~$ # Add the Kubernetes GPG Key and save it to the k
eyring
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gp
g --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg

# Add the Kubernetes repository to your APT sources list
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://p
kgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/ku
bernetes.list

# Update package lists
sudo apt update
File '/etc/apt/keyrings/kubernetes-apt-keyring.gpg' exists. Overwrite? (y/N)
y
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8
s.io/core:/stable:/v1.31/deb/ /
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelea
se
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:6 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:5 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/s
table:/v1.31/deb InRelease [1186 B]
Get:7 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/s
table:/v1.31/deb Packages [4865 B]
Fetched 6051 B in 1s (6276 B/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
142 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s
) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file h
as an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is st
ored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATI
ON section in apt-key(8) for details.
ubuntu@ip-172-31-82-192:~$
```


sudo apt-get update

sudo apt-get install -y kubelet kubeadm kubectl

sudo apt-mark hold kubelet kubeadm kubectl

```
ubuntu@ip-172-31-82-192:~$ sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:6 http://security.ubuntu.com/ubuntu noble-security InRelease
Hit:5 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb InRelease
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  conntrack cri-tools kubernetes-cni
The following NEW packages will be installed:
  conntrack cri-tools kubeadm kubectl kubelet kubernetes-cni
0 upgraded, 6 newly installed, 0 to remove and 142 not upgraded.
Need to get 87.4 MB of archives.
After this operation, 314 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 conntrack amd64 1:1.4.8-1ubuntu1 [37.9 kB]
Get:2 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb cri-tools 1.31.1-1.1 [15.7 MB]
Get:3 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubeadm 1.31.1-1.1 [11.4 MB]
Get:4 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubectl 1.31.1-1.1 [11.4 MB]
Processing triggers for man-db (2.12.0-4ubuntu2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
ubuntu@ip-172-31-82-192:~$
```

sudo systemctl enable --now kubelet

sudo apt-get install -y containerd

```
ubuntu@ip-172-31-82-192:~$ sudo systemctl enable --now kubelet
sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  runc
The following packages will be REMOVED:
  containerd.io docker-ce
The following NEW packages will be installed:
  containerd runc
0 upgraded, 2 newly installed, 2 to remove and 142 not upgraded.
Need to get 47.2 MB of archives.
After this operation, 53.1 MB disk space will be freed.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64
```

```
Setting up containerd (1.7.12-0ubuntu4.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
```

sudo mkdir -p /etc/containerd

sudo containerd config default | sudo tee /etc/containerd/config.toml

```
ubuntu@ip-172-31-82-192:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml

disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""
```

```

[proxy_plugins]

[stream_processors]

    [stream_processors."io.containerd.ocicrypt.decoder.v1.tar"]
        accepts = ["application/vnd.oci.image.layer.v1.tar+encrypted"]
        args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
        env = ["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprovider.conf"]
        path = "ctd-decoder"
        returns = "application/vnd.oci.image.layer.v1.tar"

    [stream_processors."io.containerd.ocicrypt.decoder.v1.tar.gzip"]
        accepts = ["application/vnd.oci.image.layer.v1.tar+gzip+encrypted"]
        args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
        env = ["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprovider.conf"]
        path = "ctd-decoder"
        returns = "application/vnd.oci.image.layer.v1.tar+gzip"

[timeouts]
"io.containerd.timeout.bolt.open" = "0s"
"io.containerd.timeout.metrics.shimstats" = "2s"
"io.containerd.timeout.shim.cleanup" = "5s"
"io.containerd.timeout.shim.load" = "5s"
"io.containerd.timeout.shim.shutdown" = "3s"
"io.containerd.timeout.task.state" = "2s"

[ttrpc]
address = ""
gid = 0
uid = 0
ubuntu@ip-172-31-82-192:~$

```

sudo systemctl restart containerd

sudo systemctl enable containerd

sudo systemctl status containerd

```

ubuntu@ip-172-31-82-192:~$ sudo systemctl restart containerd
sudo systemctl enable containerd
sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; p
   Active: active (running) since Thu 2024-09-26 16:19:13 UTC; 499ms ago
     Docs: https://containerd.io
   Main PID: 5387 (containerd)
      Tasks: 6
   Memory: 15.0M (peak: 15.4M)
      CPU: 78ms
   CGroup: /system.slice/containerd.service
           └─5387 /usr/bin/containerd

Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:13
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
Sep 26 16:19:13 ip-172-31-82-192 systemd[1]: Started containerd.service - c
Sep 26 16:19:13 ip-172-31-82-192 containerd[5387]: time="2024-09-26T16:19:1
lines 1-21/21 (END)

```

sudo apt-get install -y socat

```
ubuntu@ip-172-31-82-192:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 142 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat
amd64 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (15.7 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-82-192:~$ |
```

Step 6: Initialize the Kubercluster. Now perform this on Master Instance.

sudo kubeadm init --pod-network-cidr=10.244.0.0/16

```

ubuntu@ip-172-31-82-192:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=NumCPU,Mem
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING NumCPU]: the number of available CPUs 1 is less than the required 2
[WARNING Mem]: the system RAM (957 MB) is less than the minimum 1700 MB
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
[certs] Using certificateDir folder '/etc/kubernetes/pki'
[certs] Generating 'ca' certificate and key
[certs] Generating 'apiserver' certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-82-192 kubernet.es.default kubernet.es.default.svc kubernet.es.default.svc.cluster.local] and IPs [10.96.0.1 172.31.82.192]
[certs] Generating 'apiserver-kubelet-client' certificate and key
[certs] Generating 'front-proxy-ca' certificate and key
[certs] Generating 'front-proxy-client' certificate and key
[certs] Generating 'etcd/ca' certificate and key
[certs] Generating 'etcd/server' certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-82-192 localhost] and IPs [172.31.82.192 127.0.0.1 ::1]
[certs] Generating 'etcd/peer' certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-82-192 localhost] and IPs [172.31.82.192 127.0.0.1 ::1]
[certs] Generating 'apiserver-etcd-client' certificate and key
[certs] Generating 'sa' key and public key
[kubeconfig] Using kubeconfig folder '/etc/kubernetes'
[kubeconfig] Writing 'admin.conf' kubeconfig file
[kubeconfig] Writing 'super-admin.conf' kubeconfig file
[kubeconfig] Writing 'kubect.conf' kubeconfig file
[kubeconfig] Writing 'controller-manager.conf' kubeconfig file
[kubeconfig] Writing 'scheduler.conf' kubeconfig file
[etcd] Creating static Pod manifest for local etcd in '/etc/kubernetes/manifests'
[control-plane] Using manifest folder '/etc/kubernetes/manifests'
[control-plane] Creating static Pod manifest for 'kube-apiserver'
[control-plane] Creating static Pod manifest for 'kube-controller-manager'
[control-plane] Creating static Pod manifest for 'kube-scheduler'
[kubelet-start] Writing kubelet environment file with flags to file '/var/lib/kubelet/kubeadm-flags.env'
[kubelet-start] Writing kubelet configuration to file '/var/lib/kubelet/config.yaml'
[kubelet-start] Starting the kubelet
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory '/etc/kubernetes/manifests'
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.53248979s
[api-check] Waiting for a healthy API server. This can take up to 4m0s
[api-check] The API server is healthy after 11.585850126s
[upload-config] Storing the configuration used in ConfigMap 'kubeadm-config' in the 'kubernetes-system' Namespace
[kubelet] Creating a ConfigMap 'kubelet-config' in namespace kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node ip-172-31-82-192 as control-plane by adding the labels: [node-role.kubernetes.io/control-plane node.kubernetes.io/exclude-from-external-load-balancers]
[mark-control-plane] Marking the node ip-172-31-82-192 as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: jhztz.eyJ0dW83Zq4eQr
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the etcdprover controller to automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the 'cluster-info' ConfigMap in the 'kubernetes-public' namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run 'kubectl apply -f [podnetwork].yaml' with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.82.192:6443 --token jhztz.eyJ0dW83Zq4eQr \
--discovery-token-ca-cert-hash sha256:a6037b3c6680d5fdad08dfd180793d17c759d0beb046a5cc0908a9072e52d2855

```

From this command we get token and ca-

Run this command on master and

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```

ubuntu@ip-172-31-82-192:~$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-82-192:~$

```

Step 7: Now Run the command **kubectl get nodes** to see the nodes before executing Join command on nodes.

```

ubuntu@ip-172-31-82-192:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE     VERSION
ip-172-31-82-192    NotReady control-plane 7m21s    v1.31.1
ubuntu@ip-172-31-82-192:~$

```

Step 8: Now Run the following command on Node 1 and Node 2 to Join to master.

```
sudo kubeadm join <your-master-node-ip>:6443 --token <your-token>
```

```
--discovery-token-ca-cert-hash sha256:<your-ca-cert-hash>
```

kubeadm join 172.31.82.192:6443 --token jrhztc.eyi07duk03zq4eqr \
--discovery-token-ca-cert-hash
sha256:a6037b3c6608d5fdadd8dfd100793d17c759dbeb046a5ccd908a90f2e52d2055
(SLASH '\' MIGHT GIVE ERROR IF IT GIVES ERROR THEN TRY WITHOUT ERROR)

NODE1:

```
ubuntu@ip-172-31-86-84:~$ sudo kubeadm join 172.31.82.192:6443 --token jrhztc.eyi07duk03zq4eqr --discovery-token-ca-cert-hash sha256:a6037b3c6608d5fdadd8dfd100793d17c759dbeb046a5ccd908a90f2e52d2055
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.004592587s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
ubuntu@ip-172-31-86-84:~$
```

NODE2:

```
ubuntu@ip-172-31-86-8:~$ sudo kubeadm join 172.31.82.192:6443 --token jrhztc.eyi07duk03zq4eqr --discovery-token-ca-cert-hash sha256:a6037b3c6608d5fdadd8dfd100793d17c759dbeb046a5ccd908a90f2e52d2055
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.003588565s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
ubuntu@ip-172-31-86-8:~$
```

Step 9: Now Run the command on Master **kubectl get nodes** to see the nodes after executing Join command on nodes.

```
ubuntu@ip-172-31-82-192:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE     VERSION
ip-172-31-82-192    NotReady control-plane 55m     v1.31.1
ip-172-31-86-8      NotReady <none>    7m44s    v1.31.1
ip-172-31-86-84     NotReady <none>    7m5s     v1.31.1
ubuntu@ip-172-31-82-192:~$
```

Step 10: Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

```
ubuntu@ip-172-31-82-192:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
ubuntu@ip-172-31-82-192:~$
```


sudo systemctl status kubelet

```
ubuntu@ip-172-31-82-192:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Thu 2024-09-26 16:55:46 UTC; 57min ago
     Docs: https://kubernetes.io/docs/
   Main PID: 7822 (kubelet)
    Tasks: 10 (limit: 1130)
   Memory: 49.2M (peak: 72.1M)
      CPU: 44.924s
   CGroup: /system.slice/kubelet.service
            └─7822 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kube

Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: E0926 17:52:55.304654 7822 kuberuntime_>
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: rpc error: code = Unknown desc = f>
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: : unknown
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: > pod="kube-system/kube-scheduler-ip-172->
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: E0926 17:52:55.393676 7822 log.go:32] ">
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: rpc error: code = Unknown desc = f>
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: : unknown
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: > podSandboxID="af90c21fb06b79773120a5ffc>
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: E0926 17:52:55.393747 7822 kuberuntime_>
Sep 26 17:52:55 ip-172-31-82-192 kubelet[7822]: E0926 17:52:55.393869 7822 kubelet.go:1>
lines 1-23/23 (END)
```

Now Run command **kubectl get nodes -o wide** we can see Status is ready.

```
ubuntu@ip-172-31-82-192:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-82-192    Ready    control-plane   58m   v1.31.1   172.31.82.192 <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-8      Ready    <none>        10m   v1.31.1   172.31.86.8   <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-84     Ready    <none>        9m42s v1.31.1   172.31.86.84  <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
```

Now to Rename run this command

kubectl label node <node-ip> kubernetes.io/role=worker

Rename to Node 1: **kubectl label node ip-172-31-86-8 kubernetes.io/role=Node1**

Rename to Node 2: **kubectl label node ip-172-31-86-84 kubernetes.io/role=Node2**

```
kube-scheduler ip-172-31-82-192 172.31.82.192 Running 2 (11m ago) 09m
ubuntu@ip-172-31-82-192:~$ kubectl label node ip-172-31-86-8 kubernetes.io/role=Node1
^[A^[[Anode/ip-172-31-86-8 labeled
ubuntu@ip-172-31-82-192:~$ kubectl label node ip-172-31-86-84 kubernetes.io/role=Node2
node/ip-172-31-86-84 labeled
ubuntu@ip-172-31-82-192:~$ |
```

Step 11: Run command **kubectl get nodes -o wide** . And Hence we can see we have Successfully connected Node 1 and Node 2 to the Master.

```
ubuntu@ip-172-31-82-192:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-82-192    Ready    control-plane   83m   v1.31.1   172.31.82.192 <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-8      Ready    Node1        35m   v1.31.1   172.31.86.8   <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-84     Ready    Node2        34m   v1.31.1   172.31.86.84  <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
^[Aubuntu@ip-172-31-82-192:~$ kubectl get nodes -o wide|
```

Or run **kubectl get nodes**

```
^[Aubuntu@ip-172-31-82-192:kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-82-192    Ready    control-plane   85m   v1.31.1   172.31.82.192 <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-8      Ready    Node1        37m   v1.31.1   172.31.86.8   <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ip-172-31-86-84     Ready    Node2        36m   v1.31.1   172.31.86.84  <none>         Ubuntu 24.04 LTS    6.8.0-1012-aws    containerd://1.7.12
ubuntu@ip-172-31-82-192:~$ |
```

Conclusion:

In this Advanced DevOps Lab experiment, we began by setting up three EC2 Ubuntu instances on AWS, designating one as the Master node and the others as Worker nodes. We then installed Docker and Kubernetes on all instances, ensuring Docker was properly configured. The Kubernetes cluster was initialized on the Master node, and the Flannel networking plugin was applied to facilitate communication between nodes. Finally, we joined the Worker nodes to the cluster using the provided token and hash, resulting in a fully operational Kubernetes cluster ready for managing and scaling containerized applications.