

## Advanced DevOps Lab

### Experiment 4

**Aim:** To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

#### Theory:

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications. In fact, Kubernetes has established itself as the defacto standard for container orchestration and is the flagship project of the Cloud Native Computing Foundation (CNCF), backed by key players like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

#### Kubernetes Deployment

A Kubernetes Deployment is used to tell Kubernetes how to create or modify instances of the pods that hold a containerized application. Deployments can scale the number of replica pods, enable the rollout of updated code in a controlled manner, or roll back to an earlier deployment version if necessary.

#### Steps:

##### 1. A] Creation Of EC-2 instance

→ Launch an AWS EC2 instance named AdityaRaorane using an AWS Linux AMI. Configure the Security Group's Inbound Rules to allow SSH access, then choose the t2.micro instance type.

Name

exp\_4\_adv\_devops [Add additional tags](#)

▼ **Application and OS Images (Amazon Machine Image)** [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Q Search our full catalog including 1000s of application and OS images

Recents Quick Start

Amazon Linux macOS Ubuntu Windows Red Hat SUSE Linux

[Browse more AMIs](#)

Including AMIs from AWS, Marketplace and the Community

Number of instances [See more](#)

1

[Software Image \(AMI\)](#)

Amazon Linux 2023 AMI 2023.5.0  
ami-0182f373e66f89c85

[Virtual server type \(instance type\)](#)

t2.micro

[Firewall \(security group\)](#)

New security group

[Storage \(volumes\)](#)

1 volume(s) - 8 GiB

**Free tier:** In your first year includes 750 hours of t2.micro in the Regions in t2.micro is unavailable) in usage on free tier AMIs

Cancel



```
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-29-131 docker]$ cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts":
["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
{
"exec-opts":
["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
[ec2-user@ip-172-31-29-131 docker]$
```

After this run the following command to enable and start docker and also to load the daemon.json file.

```
→ sudo systemctl enable docker          → sudo systemctl daemon-reload
→ sudo systemctl restart docker         → docker -v
```

```
[ec2-user@ip-172-31-29-131 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-29-131 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-29-131 docker]$ sudo systemctl restart docker
[ec2-user@ip-172-31-29-131 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-29-131 docker]$
```

### 3. Install Kubernetes →

a) SELinux needs to be disabled before configuring kubelet Run the following command

```
→ sudo setenforce 0
```

```
→ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

b) We are adding kubernetes using the repository whose command is given below.

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
```

```
enabled=1
```

```
gpgcheck=1
```

```

gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.x
ml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF

```

```

[ec2-user@ip-172-31-29-131 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-29-131 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo/' /etc/selinux/config
[kubernetes]-172-31-29-131 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.x
ml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cniomd.x
EOFkey exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.x
ml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[ec2-user@ip-172-31-29-131 docker]$

```

c) After that Run following command to make the updation and also to install kubelet, kubeadm, kubectl: → sudo yum update

→sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

```

[ec2-user@ip-172-31-29-131 docker]$ sudo yum update
Kubernetes               113 kB/s | 17 kB      00:00
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-29-131 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:01:14 ago on Sat Sep 14 13:55:34 2024.
Dependencies resolved.
=====
Package                Arch    Version                Repository              Size
=====
Installing:
kubeadm                x86_64  1.30.5-150500.1.1     kubernetes              10 M
kubectl                x86_64  1.30.5-150500.1.1     kubernetes              10 M
kubelet                x86_64  1.30.5-150500.1.1     kubernetes              17 M
Installing dependencies:
conntrack-tools        x86_64  1.4.6-2.amzn2023.0.2   amazonlinux             208 k
cri-tools              x86_64  1.30.1-150500.1.1     kubernetes              8.6 M
kubernetes-cni         x86_64  1.4.0-150500.1.1     kubernetes              6.7 M
libnetfilter_cthelper  x86_64  1.0.0-21.amzn2023.0.2   amazonlinux             24 k
libnetfilter_cttimeout x86_64  1.0.0-19.amzn2023.0.2   amazonlinux             24 k
libnetfilter_queue     x86_64  1.0.5-2.amzn2023.0.2   amazonlinux             30 k
Transaction Summary
=====
Install  9 Packages
Total download size: 53 M

```

d) After installing Kubernetes, we need to configure internet options to allow bridging.

1. sudo swapoff -a

2. echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf

3. sudo sysctl -p

```

[ec2-user@ip-172-31-29-131 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-29-131 docker]$

```

#### 4. Initialize the Kubecluster

→ a] Initializes a Kubernetes cluster with kubeadm, sets up the pod network CIDR to 10.244.0.0/16 for network communication, and ignores preflight checks for CPU and memory requirements.

*sudo kubeadm init --ignore-preflight-errors=NumCPU,Mem*

```
[root@ip-172-31-26-1 docker]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@ip-172-31-26-1 docker]# kubeadm init
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
        [WARNING FileExisting-socat]: socat not found in system path
        [WARNING FileExisting-tc]: tc not found in system path
error execution phase preflight: [preflight] Some fatal errors occurred:
        [ERROR NumCPU]: the number of available CPUs 1 is less than the required 2
        [ERROR Mem]: the system RAM (949 MB) is less than the minimum 1700 MB
[preflight] If you know what you are doing, you can make a check non-fatal with `--ignore-preflight-errors=...`
To see the stack trace of this error execute with --v=5 or higher
[root@ip-172-31-26-1 docker]# kubeadm init --ignore-preflight-errors=NumCPU,Mem
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
        [WARNING NumCPU]: the number of available CPUs 1 is less than the required 2
        [WARNING Mem]: the system RAM (949 MB) is less than the minimum 1700 MB
```

b] copy the token and save for future use .

```
kubeadm join 172.31.26.1:6443 --token s9t4e8.ygz2u9mm5okam1vq \
--discovery-token-ca-cert-hash
sha256:45eca28e5c3fee977bf721e09d48012e10734c339937e7742b68c3bfd538c26c
```

c] Copy the mkdir and chown commands from the top and execute them

```
→ mkdir -p $HOME/.kube
→ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
→ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
kubeadm join 172.31.26.1:6443 --token s9t4e8.ygz2u9mm5okam1vq \
--discovery-token-ca-cert-hash sha256:45eca28e5c3fee977bf721e09d48012e10734c339937e7742b68c3bfd538c26c
```

d] Then, add a common networking plugin called flannel as mentioned in the code.

kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[root@ip-172-31-26-1 docker]# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

**5. Deploying an NGINX Server on Your Kubernetes Cluster** → a] Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply deployment using this following command:

*kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml*

```
bash: https://k8s.io/examples/pods/simple-pod.yaml: No such file or directory
[root@ip-172-31-26-1 docker]# kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
[root@ip-172-31-26-1 docker]#
```

b] Then use **kubectl get pods** to check whether the pod gets created or not.

```
[root@ip-172-31-26-1 docker]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Pending	0	54s

c)To convert state from pending to running use following command:

**kubectl describe pod nginx** This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted

```
[ec2-user@ip-172-31-26-1 ~]$ sudo nano nginx-deployment.yaml
[ec2-user@ip-172-31-26-1 ~]$ kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx created
[ec2-user@ip-172-31-26-1 ~]$ kubectl get pods
No resources found in default namespace.
[ec2-user@ip-172-31-26-1 ~]$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx	0/1	1	0	78s

```
[ec2-user@ip-172-31-26-1 ~]$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-7769f8f85b-wlmzm	0/1	Pending	0	80s

```
[ec2-user@ip-172-31-26-1 ~]$ kubectl describe deployment nginx
Name:
Namespace:
CreationTimestamp:
Labels:
Annotations:
Selector:
```

nginx	default	Sat, 14 Sep 2024 19:16:31 +0000	<none>	deployment.kubernetes.io/revision: 1	app=nginx
-------	---------	---------------------------------	--------	--------------------------------------	-----------

**FOR ERROR (paste this in nginx-deployment.ymal & follow above commands)**

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx

spec:

replicas: 1

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

tolerations:

- key: "node-role.kubernetes.io/control-plane"

operator: "Exists"

effect: "NoSchedule"

containers:

- name: nginx

image: nginx

ports:  
- containerPort: 80

```
Available      False      MinimumReplicasUnavailable
Progressing    True       ReplicaSetUpdated
OldReplicaSets: <none>
NewReplicaSet: nginx-7769f8f85b (1/1 replicas created)
Events:
-----
Type          Reason            Age          From          Message
-----
Normal        ScalingReplicaSet  98s         deployment-controller  Scaled up replica set nginx-7769f8f85b to 1
[ec2-user@ip-172-31-26-1 ~]$ kubectl get events
LAST SEEN   TYPE      REASON              OBJECT                                          MESSAGE
10m         Normal    NodeHasSufficientMemory  node/ip-172-31-26-1.ec2.internal             Node ip-172-31-26-1.ec2.internal status is now: NodeHasSufficientMemory
10m         Normal    NodeHasNoDiskPressure   node/ip-172-31-26-1.ec2.internal             Node ip-172-31-26-1.ec2.internal status is now: NodeHasNoDiskPressure
10m         Normal    NodeHasSufficientPID     node/ip-172-31-26-1.ec2.internal             Node ip-172-31-26-1.ec2.internal status is now: NodeHasSufficientPID
10m         Normal    NodeAllocatableEnforced  node/ip-172-31-26-1.ec2.internal             Updated Node Allocatable limit across pods
```

6. Now check pod status is is running

```
NAME                                STATUS    ROLES    AGE    VERSION
ip-172-31-26-1.ec2.internal        Ready    control-plane  83s    v1.31.1
[ec2-user@ip-172-31-26-1 ~]$
```

7. Lastly, mention the port you want to host. Here i have used localhost 8081 then check it.

**kubectl port-forward nginx 8081:80**

```
[ec2-user@ip-172-31-73-36 docker]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

8. Verify your deployment Open up a new terminal and ssh to your EC2 instance. Then, use this curl command to check if the Nginx server is running.

**curl --head http://127.0.0.1:8080**

If the response is 200 OK and you can see the Nginx server name, your deployment was successful. We have successfully deployed our Nginx server on our EC2 instance.

**Conclusion:** We successfully set up a Kubernetes cluster on AWS EC2, addressing issues related to component setup and residual configurations. We ensured proper cleanup of previous Kubernetes files and mounts, verified the kubelet service, and applied Flannel for networking. Finally, we resolved connectivity issues, and after a thorough review of logs and configuration, we deployed and exposed an NGINX server using Kubernetes services, preparing the cluster for efficient traffic management and scaling.